



RQ1 performance summary: representation effect and baseline context

The primary LLM comparison uses 178 complete pairs; supporting baselines use 200 fixed five-fold OOF predictions.

(a) Primary LLM representation experiment



(b) Supporting tabular baselines



Primary paired accuracy effect: structured – text = +0.079 | Exact McNemar p = 0.0436

The panels use a common 0–1 scale but different analysis populations. No formal statistical test directly compares the LLM conditions with the tabular baselines.

Baseline performance shows that the locked 46-feature table contains strong within-dataset signal; it does not establish external generalisation.