


Is LLM-model alignment specific to the same observation?



Grey intervals show the central 95% of 50,000 class-stratified strict-mismatch randomisations. Diamonds show observed same-sample alignment. The null preserves true class and global feature preferences.