


Effect of input representation on LLM-model attribution alignment

Deterministic text higher Structured higher



Points are paired mean differences across 178 samples; bars are 95% paired percentile-bootstrap confidence intervals (20,000 resamples). P-values use 100,000 paired sign-flip randomisations and Holm family-wise-error correction across four hypotheses.